

HW3 - Web Servers and Synchronization

Due Date: 6.7.2026

TAs in charge: Shachar Asher Cohen – cohenshachar@campus.technion.ac.il

George Nasser - georgenasser@campus.technion.ac.il

Important: The Q&A for the exercise will take place at a public forum Piazza only.

Critical updates about the HW will be published in pinned notes in the piazza forum. These notes are mandatory and it is your responsibility to be updated. A number of guidelines to use the forum:

- Read previous Q&A carefully before asking the question; repeated questions will probably go without answers.
- Be polite, remember that course staff does this as a service for the students.
- You're not allowed to post any kind of solution and/or source code in the forum as a hint for other students; In case you feel that you have to discuss such a matter, please come to the reception hour.
- When posting questions regarding this assignment, put them in the "hw3-wet" folder. Pay attention that in the dry part there are questions about the wet part, so it is recommended to review these questions before, and answer them gradually while working on the wet part.

Only the TA in charge, can authorize postponements. In case you need a postponement, please fill out the attached form: [Spring2026 - HW3 postponement request – מלא את הטופס](#)

Introduction

"This assignment PDF is so long OMG!" well it is mostly going to be explanations, the tl;dr of what you are required to do is written as "Task N:..." throughout the assignment. If you find it less confusing to start by reading the tasks and looking for the relevant explanations

afterwards, go for it.

In this assignment, you will develop a **working multi-threaded web server**. The server will maintain a **log of handled jobs** ("the server's book") and allow **multiple concurrent clients** to **read** or **add entries** to this log.

To simplify the project, we provide a basic single-threaded web server that already:

- Handles both **GET** and **POST** requests
- Includes built-in response composition
- Abstracts away low-level networking and HTTP handling
- A stub log implementation and more necessary structs for quality of life.

Your task is:

- 1) to make this server **multi-threaded** for efficiency.
- 2) add statistic collection as will be described later.
- 3) add logging **functionality**.
- 4) add a logging debugging feature

HTTP Background

This section is a brief overview of the **HTTP protocol**, which is already handled for you in the starter code.

- Web browsers and servers communicate via **HTTP** (Hypertext Transfer Protocol), a text-based request/response protocol.
- The browser opens a connection, makes a request (e.g., for a file), the server responds, and the connection is closed.
- **GET** requests are typically used to retrieve static content (such as HTML files) from the server. Each file has a unique name known as a URL (Universal Resource

Locator).

For example, the URL: www.cs.technion.ac.il:80/index.html identifies an HTML file called "/index.html" on Internet host "www.cs.technion.ac.il" that is managed by a web server listening on port 80.

- URLs can also reference **dynamic content**, like CGI scripts, using the ? and & characters to pass arguments, e.g. A request like: "GET /req.cgi?user=alex&id=5" is asking to run req.cgi with the parameters user = alex and id = 5 and return the output. (If you're interested, you can read more about CGI in the attached link- [what is cgi programming - Stack Overflow](#)).
- **POST** requests are used to submit data to the server, often resulting in a change in server state or side effects (e.g., updating a database or creating a resource).
- **In this assignment:**
 - **GET request** behaves similarly to a traditional GET request, with one key difference - It also **appends an entry** representing the request itself to the book **before** returning it.
 - **POST request** retrieves the current contents of the **server's book**, which is a log of all requests previously made to the server.

HTTP Request Structure:	HTTP Response Structure:
consists of a request line, followed by zero or more request headers, and finally an empty text line. A request line has the form: [method] [uri] [version]. The method is usually GET (but may be other things, such as POST, OPTIONS, PUT or whatever the server knows how to handle).	consists of a response line, zero or more response headers, an empty text line, and the response body. A response line has the form [version] [status] [message]. The status is a three-digit positive integer that indicates the state of the request; some common states are 200 for "OK", 403 for "Forbidden", and 404 for "Not found". Two important lines in the header are "Content-Type", which tells the client the MIME type

	of the content in the response body (e.g., html or gif) and "Content-Length", which indicates its size in bytes.
[method] [uri] [version] [header1]: [value1] [header2]: [value2] [optional body for POST requests]	[version] [status code] [message] Content-Type: text/html Content-Length: 123 [body content]

If you would like to see the HTTP protocol in action, you can connect to any web server using telnet. For example, run `telnet www.cs.technion.ac.il 80` and then type:

```
GET / HTTP/1.1
host: www.cs.technion.ac.il
[empty line - simply press enter twice after previous line]
```

You will then see the raw HTML text for that web page!

You do **not need to modify** any of the code related to HTTP parsing or network connections. The starter code fully handles these aspects for you. However, understanding this background will help you reason about how requests are structured and processed internally and how to complete the assignment.

Basic Web Server -Single Thread, Single Port

The code for the web server is available on the course website. You can compile the files by simply typing `make`. Compile and run this basic web server before making any changes to it! `make clean` removes .o files and executables and lets you do a clean build.

When you run this basic web server, you need to specify the port number that it will listen on; ports below number 1024 are *reserved* (see the list here) so you should specify port numbers that are greater than 1023 to avoid this reserved range; the max is 65535. Be wary: if running on a shared machine, you could conflict with others and thus have your server fail to bind to the desired port. If this happens, try a different number!

You can run the server by running:

```
./server [desired_port]
```

In addition, we have provided you with a very primitive HTTP client. This client can request a file from any HTTP server by running:

```
./client [hostname] [port] [filename] [method]
```

The server supports the methods GET and POST.

When you connect to the server, make sure that you specify this same port. For example, assume that you started the server on csl3.cs.technion.ac.il and used port number 8003, by running:

```
./server 8003
```

The server can respond with any html file you put in the public directory (automatically generated by make). Let's assume you copied there your favorite html file - favorite.html.

To view this file using the provided client (running on the same or a different machine), run:

```
./client csl3.cs.technion.ac.il 8003 favorite.html
```

To view this file from a web browser (running on the same or a different machine), use the url:

```
csl3.cs.technion.ac.il:8003/favorite.html
```

If you run the client and web server on the same machine, you can just use the hostname localhost as a convenience, e.g., localhost:8003/favorite.html.

For your convenience we added the basic home.html file for you to experiment with.

Notes

1. On some browsers you might need to use an incognito tab.

2. If you use csl3, your client/browser must be connected to the Technion's network. If you don't have access to the Technion's network (or just don't want to use csl3), you can always run the server locally on your machine and access it with localhost.

To make the homework a bit easier, the web server is very minimal, consisting of only a few hundred lines of C code. As a result, the server is limited in its functionality; it does not handle any HTTP requests other than **GET** and **POST**, understands only a few content types, and supports only the `QUERY_STRING` environment variable for CGI programs. This web server is also not very robust; for example, if a web client closes its connection to the server, it may trip an assertion in the server causing it to exit. We do not expect you to fix these problems (though you can, if you like, you know, for fun).

Helper functions are provided to simplify error checking. A wrapper calls the desired function and immediately terminates if an error occurs. The wrappers are found in the file `sege1.h`); more about this below. **One should always check error codes, even if all you do in response is exit**; dropping errors silently is BAD C PROGRAMMING and should be avoided at all costs.

Part 1: Multi-threaded Web Server — Detailed Explanation and Guidance

Background:

The provided starter server is single-threaded — it can only handle one HTTP request at a time. This design leads to performance bottlenecks, especially when:

- A request takes a long time to complete (e.g., running a CGI script).
- Disk I/O or network latency causes delays in responding.

Furthermore, It only handles heavy TCP connection for the HTTP files and not quick telemetry output.

Since the server blocks on one request, all other incoming clients must wait, even for trivial requests. This results in poor scalability and user experience.

The simplest multithreading strategy is to create **a new thread for every incoming request**. This ensures:

- Lightweight/fast requests don't get stuck behind slow ones.

- Threads blocked on I/O don't hold up the rest of the server.

However, this approach has downsides:

- **Thread creation overhead** on each request.
- Potential performance degradation under high load.

Task 1: Implement a Fixed-Size Thread Pool Producer-Consumer Server

Overview

You will modify the provided basic web server to handle multiple HTTP requests concurrently by implementing:

- One master thread (the producer) that accepts incoming client connections.
- A fixed-size pool of worker threads (the consumers) that process these requests in parallel.

This design is often called a thread pool and is more efficient than creating a new thread for every request because:

- Thread creation overhead is avoided on each request.
- Thread management is centralized and controlled.
- It allows better control over server resource usage and load.

How It Works

1. **Master Thread (Producer):**
 - Listens on the server socket for new client connections.
 - When a connection is accepted, it adds the connection to a bounded FIFO queue.
 - If the queue is full, the master thread blocks (waits) until space becomes available.
 - Signals worker threads that new work is ready.
2. **Worker Threads (Consumers):**
 - Created once at server startup; the pool size is fixed.
 - Each worker thread waits for work to appear in the queue.
 - When awakened, a worker picks the oldest request from the queue (FIFO).

- Processes the request (handle GET or POST) – provided function see request.h/c.
- Returns to waiting for the next request.

Synchronization Requirements

Because multiple threads (master + workers) access shared resources (the request queue, the server's log), synchronization is crucial.

- Avoid race conditions.
- Use condition variables to avoid busy-waiting:
 - Workers wait when the queue is empty.
 - Master waits when the queue is full.
- No spin-locks or busy waiting allowed – will be penalized heavily. (i.e., no looping while checking conditions without sleeping).

Handling GET and POST Requests with Reader-Writer Synchronization:

The server's will be logging valid requests it has handled. A shared data structure called Server_Log (see log.h/c) maintains a log of all handled requests.

- **GET requests:**
 - Are considered writers to the book — they append an entry to the log.
 - Require exclusive access to the book.
- **POST requests:**
 - Are considered readers — they read and return the entire log.
 - Multiple POST requests (readers) can access the book simultaneously.

Task 2: You must implement the log (see log.c/h) with a reader-writer synchronization scheme with writer priority:

- If readers and writers are waiting simultaneously, give precedence to writers.
- This ensures writers don't starve behind a flood of readers.

Edge Cases and Important Notes

- **Queue Size Limit:** The bounded queue has a fixed capacity (queue_size).
 - When full, the master thread must block, so new connections wait until there is space.
- **Queue Management:**

- You may implement one queue for pending requests.
- Optionally, you can maintain a second structure for requests currently being processed to enforce queue size constraints (sum of pending + active $\leq$ queue_size).
- All request count for queue size regardless of type.
- Worker threads should never block other threads while processing their assigned request. This means once a worker takes a request off the queue, it releases the lock so others can continue.
- The basic server forks a process for CGI in requestPrepareDynamic
 - Your multi-threaded server must not modify this CGI forking logic.
 - The server currently waits for each CGI child process before continuing.
- FIFO order of requests must be maintained when worker threads pick requests from the queue.
- Proper error checking and robust synchronization (no race conditions, no deadlocks) are required.

Implementation Hints

- Use pthread_mutex_t for mutexes and pthread_cond_t for condition variables.
- Initialize your thread pool and synchronization primitives at server startup.
- Use functions provided in segel.c for error-checked system calls and synchronization primitives.
- Implement the reader-writer lock for the server's book log separately, ensuring:
 - Multiple concurrent POST readers.
 - Exclusive GET writers with priority.

Part 2: Usage Statistics – Multi-Port Server

So far our server should be capable of handling multiple TCP requests in a FIFO manner using a pooled thread architecture. To help the admins of the server monitor the server for backlogs, your server will support statistics collection and a UDP port connection alongside the TCP connection. This UDP port will get pinged by admins for specific thread statistics.

You will need to modify your web server to collect a variety of statistics, some of which statistics will be gathered on a per-request (job) basis and some on a per-thread basis. Job related statistics will be returned alongside the response in the http header (we have already made place-holders in the basic web server code for these headers (Check out how the server sends "Content-length" and "Content-type" and what the client does with this information); while thread and system related statistic will be returned in response to UDP pings.

JOB STATISTICS

For each request, you will record the following times or durations at the granularity of milliseconds. You may find [`gettimeofday\(\)`](#) useful for gathering these statistics.

Task 3: collect the statistics according to the instructions:

- **Stat-req-arrival:** The arrival time, as first seen by the master thread

format:

```
Stat-Req-Arrival:: %ld.%06ld
```

- **Stat-req-dispatch:** The time at which a worker thread lifted the job from the queue

format:

```
Stat-Req-Dispatch:: %ld.%06ld
```

- **Stat-Log-Arrival:** The time at which a request arrived at the logger

format:

```
Stat-Log-Arrival:: %ld.%06ld
```

- **Stat-Log-Dispatch:** The time at which log finish handling said request

format:

```
Stat-Log-Dispatch:: %ld.%06ld
```

Note: more instructions on log arrival and log dispatch timing in task 5.

THREAD STATISTICS

- **Stat-thread-id:** The id of the responding thread where worker threads are distributed id in numbers from 1 to N.

format:

```
Stat-Thread-Id:: %d
```

- **Stat-thread-count:** The total number of http requests this thread has handled

format:

```
Stat-Thread-Count:: %d
```

- **Stat-thread-static:** The total number of static requests this thread has handled

format:

```
Stat-Thread-Static:: %d
```

- **Stat-thread-dynamic:** The total number of dynamic requests this thread has handled

format:

```
Stat-Thread-Dynamic:: %d
```

- **Stat-thread-post:** The total number of post requests this thread has handled

format:

```
Stat-Thread-Post:: %d
```

For a UDP ping handled by thread <id>, your web server will return that thread's statistics.

UDP pings are not considered jobs in the queue. They should run before and job is handled given the opportunity – First we respond to UDP and only then we handle TCP in the worker thread.

Task 4: Support UDP channel and response. UDP pings should be collected by the Master Thread concurrently with the TCP requests. UDP pings hold an <id> of the thread they want to read and should be registered to that thread to handle. Assume multiple UDP requests can be handed to the same thread.

Your code should follow the following logic:

- ❑ All requests (including errors) increment the request counter.
- ❑ valid static requests increment the static counter.
- ❑ valid dynamic requests increment the dynamic counter
- ❑ valid post requests increment the post counter
- ❑ 501, 403 and 404 errors do not increment the dynamic/static/post counters

*Note that some of the implementation should happen in `handle_request`.

Server Log:

The log must contain a record of all valid request statistics handled by the server (the records being the statistics gathered at handling both thread-statistics and then job-statistics. The Log is required to be returned as a chained string delimited between entries by “#” of all the requests and thread statistics that happened up to the post request.

- GET adds an entry to the book; POST reads and returns the book contents.

In addition the log should have a sleep inside its **writing and reading critical sections** with variable time, when the server is launched with a debug mode flag and the time of sleep. **Writing or reading** should activate a sleep with the specific time at each **reading\writing** action in the log.

Task 5: implement the debugging sleep inside log critical section. **Timing Note:** The debug sleep is meant to simulate slow disk I/O in a uniform way. Therefore, the thread must execute the sleep *while holding* the reader or writer lock (inside the critical section). Stat-Log-Arrival must be recorded before requesting the lock, and Stat-Log-Dispatch must be recorded after the lock is released. This means your total log processing time will naturally include the sleep duration.

Program Specifications

Your C program must be invoked exactly as follows:

`./server [tcp_portnum] [udp_portnum] [threads] [queue_size][debug_sleep_time]`

The command line arguments to your web server are to be interpreted as follows:

- **Tcp_portnum**: the port number that the web server should listen on for tcp requests; the basic web server already handles this argument. Must be above 1024
- **Udp_portnum**: the port number that the web server should listen on for udp pings; the basic web server does not handles this argument. Must be above 1024 and different than Tcp_portnum.
- **threads**: the number of worker threads that should be created within the web server. Must be a positive integer.
- **queue_size**: the number of request connections that can be accepted at one time. Must be a positive integer. Note that it is not an error for more or less threads to be created than buffer space.
- **Debug_sleep_time**: Sets the server for debugging. Any number above 0 should set the server to debug mode where **writing and readers** to the log blocks for sleep_time seconds in their critical areas. Numbers 0 and below should mean the logger does not add a deliberate blocking mechanism **on log writing and reading**. **Can be a floating number**

For example, if you run your program as

```
./server 5003 5004 8 16 0
```

then your web server will listen for tcp on port 5003, udp on 5004, create 8 worker threads for handling http requests, allocate a 16 buffers queue for connections that are currently in progress or waiting, and do not sleep on log writing.

Hints

We recommend understanding how the code that we gave you works. We provide the following .c files:

- **server.c**: Contains main() for the basic web server.
- **request.c**: Performs most of the work for handling requests in the basic web server.

- **log.c:** A stub of a log for you to implement your real logging logic connected to request.c already so see that you work correctly with what is given (you are allowed to change almost everything if you see fit to do so, in the end your server will be compiled and executed).
- **segel.c:** Contains wrapper functions for the system calls invoked by the basic web server and client. The convention is to capitalize the first letter of each routine. Feel free to add to this file as you use new libraries or system calls. You will also find a corresponding header (.h) file that should be included by all of your C files that use the routines defined here.
- **client.c:** Contains main() and the support routines for the very simple web client. To test your server, you may want to change this code so that it can send simultaneous requests to your server. At a minimum, you will want to run multiple copies of this client.
- **output.c:** Code for a CGI program that repeatedly sleeps for a random amount of time. You may find that having a CGI program that takes awhile to complete is useful for testing your server. The documentation for how to use this program is in the source file.
- You might need to add data structures, create your own implementations and add them to your submission.

We also provide you with a sample Makefile that creates server, client, and output.cgi. You can type "make" to create all of these programs. You can type "make clean" to remove the object files and the executables. You can type "make server" to create just the server program, etc. If you create new files, you will need to add them to the Makefile. **You are allowed to program only in C, and include only the headers included in segel.h.**

The best way to learn about the code is to compile and run it.

Run the server we gave you with your preferred web browser, run this server with the client code we gave you. You can even have the client code we gave you contact any other server. Make small changes to the server code (e.g., have it print out more debugging information) to see if you understand how it works.

We anticipate that you will find the following routines useful for creating and synchronizing threads: *pthread_create*, *pthread_mutex_init*, *pthread_mutex_lock*, *pthread_mutex_unlock*, *pthread_cond_init*, *pthread_cond_wait*, *pthread_cond_signal*.

Example Output

Attached are some screenshots of our basic webserver for GET(static) and POST requests for you to compare and see that you are ready to start.

```
shachar@Shachar: /mnt/c/Users/shach/Desktop/server$ ./server 8000
```

```
shachar@Shachar:/mnt/c/Users/shach/Desktop/server$ ./client localhost 8000 home.html GET
Connected to server. clientfd = 3
Request:
GET home.html HTTP/1.1
host: Shachar

Header: HTTP/1.0 200 OK
Header: Server: OS-HW3 Web Server
Header: Content-Length: 293
Length = 293
Header: Content-Type: text/html
Header: Stat-Req-Arrival:: 93918567743552.140520697862204
Header: Stat-Req-Dispatch:: 1776.140729487092761
Header: Stat-Log-Arrival:: 1.000001
Header: Stat-Log-Dispatch:: 1.000001
Header: Stat-Thread-Id:: 0
Header: Stat-Thread-Count:: 0
Header: Stat-Thread-Static:: 0
Header: Stat-Thread-Dynamic:: 0
Header: Stat-Thread-Post:: 0
<html>

<head>
  <title>OS-HW3 Test Web Page</title>
</head>

<body>

<h2> OS-HW3 Test Web Page</h2>

<p> Test web page: simply awesome.</p>

<p> Click <a href="https://www.youtube.com/watch?v=dQw4w9WgXcQ"> here</a> for something
```



```
shachar@Shachar:/mnt/c/Users/shach/Desktop/server$ ./client localhost 8000 home.html POST
Connected to server. clientfd = 3
Request:
POST home.html HTTP/1.1
host: Shachar

Header: HTTP/1.0 200 OK
Header: Server: OS-HW3 Web Server
Header: Content-Length: 24
Length = 24
Header: Content-Type: text/plain
Header: Stat-Req-Arrival:: 93918567743552.140520697862204
Header: Stat-Req-Dispatch:: 1776.140729487092761
Header: Stat-Log-Arrival:: 1.000001
Header: Stat-Log-Dispatch:: 1.000001
Header: Stat-Thread-Id:: 0
Header: Stat-Thread-Count:: 0
Header: Stat-Thread-Static:: 0
Header: Stat-Thread-Dynamic:: 0
Header: Stat-Thread-Post:: 0
Log is not implemented.
```

```
shachar@Shachar:/mnt/c/Users/shach/Desktop/server$ ./client localhost 8000 home.html FAKE
Connected to server. clientfd = 3
Request:
FAKE home.html HTTP/1.1
host: Shachar

Header: HTTP/1.0 501 Not Implemented
Header: Content-Type: text/html
Header: Content-Length: 157
Length = 157
Header: Stat-Req-Arrival:: 93918567743552.140520697862204
Header: Stat-Req-Dispatch:: 1776.140729487092761
Header: Stat-Log-Arrival:: 140729487306752.1103823372288
Header: Stat-Log-Dispatch:: 2.529267711
Header: Stat-Thread-Id:: 0
Header: Stat-Thread-Count:: 0
Header: Stat-Thread-Static:: 0
Header: Stat-Thread-Dynamic:: 0
Header: Stat-Thread-Post:: 0
<html><title>OS-HW3 Error</title><body bgcolor="ffffff">
501: Not Implemented
<p>OS-HW3 Server does not implement this method: FAKE
<hr>OS-HW3 Web Server
```

```
shachar@Shachar:/mnt/c/Users/shach/Desktop/server$ ./client localhost 8000 bad_file.html GET
Connected to server. clientfd = 3
Request:
GET bad_file.html HTTP/1.1
host: Shachar

Header: HTTP/1.0 404 Not found
Header: Content-Type: text/html
Header: Content-Length: 163
Length = 163
Header: Stat-Req-Arrival:: 93918567743552.140520697862204
Header: Stat-Req-Dispatch:: 1776.140729487092761
Header: Stat-Log-Arrival:: 140729487306752.1103823372288
Header: Stat-Log-Dispatch:: 2.529267711
Header: Stat-Thread-Id:: 0
Header: Stat-Thread-Count:: 0
Header: Stat-Thread-Static:: 0
Header: Stat-Thread-Dynamic:: 0
Header: Stat-Thread-Post:: 0
<html><title>OS-HW3 Error</title><body bgcolor="ffffff">
404: Not found
<p>OS-HW3 Server could not find this file: ./public/bad_file.html
```

Submission

You should create a zip file containing

- All of your server source files (*.c and *.h)
- a Makefile
- A file named submitters.txt which includes the ID, name and email of the participating students. The following format should be used:

```
Linus Torvalds linus@gmail.com 234567890
```

```
Ken Thompson ken@belllabs.com 345678901
```

Important Note 1: Make sure your code compiles!

Important Note 2: In a different guide we will go through debugging and submission more in depth.

Important Note 3: when you submit, keep your confirmation code and a copy of the file(s), in case of technical failure. Your confirmation code is the only valid proof that you submitted your assignment when you did. Good luck!

-- Course staff

